

e1000网卡收包callstack分析

TXT

```
/**
 * netif_receive_skb_list - process many receive buffers from
 * network
 * @head: list of skbs to process.
 *
 * Since return value of netif_receive_skb() is normally
 * ignored, and
 * wouldn't be meaningful for a list, this function returns
 * void.
 *
 * This function may only be called from softirq context and
 * interrupts
 * should be enabled.
 */
```

void netif_receive_skb_list(struct list_head *head): **drivers/net** 实现各种硬件设备的驱动，收包后从硬件获取buffer，驱动代码中调用 **netif_receive_skb_list** 进入内核网络系统，进行接下来的处理

驱动向内核注册softirq，里面包含回调函数。驱动收到数据触发中断，kernel读取

```

__netif_receive_skb_list_ptype(struct net_device * orig_dev,
struct packet_type * pt_prev, struct list_head * head)
(/data00/codes/linux/net/core/dev.c:5533)
__netif_receive_skb_list_core(struct list_head * head, bool
pfmemalloc) (/data00/codes/linux/net/core/dev.c:5582)
__netif_receive_skb_list(struct list_head * head)
(/data00/codes/linux/net/core/dev.c:5634)
netif_receive_skb_list_internal(struct list_head * head)
(/data00/codes/linux/net/core/dev.c:5725)
gro_normal_list(struct napi_struct * napi)
(/data00/codes/linux/include/net/gro.h:433)
gro_normal_list(struct napi_struct * napi)
(/data00/codes/linux/include/net/gro.h:429)
napi_complete_done(struct napi_struct * n, int work_done)
(/data00/codes/linux/net/core/dev.c:6065)
e1000_clean(struct napi_struct * napi, int budget)
(/data00/codes/linux/drivers/net/ethernet/intel/e1000/e1000_main.c:3811)
__napi_poll(struct napi_struct * n, bool * repoll)
(/data00/codes/linux/net/core/dev.c:6496)
napi_poll(struct list_head * repoll, struct napi_struct * n)
(/data00/codes/linux/net/core/dev.c:6563)
net_rx_action(struct softirq_action * h)
(/data00/codes/linux/net/core/dev.c:6696)
__do_softirq() (/data00/codes/linux/kernel/softirq.c:571)
do_softirq() (/data00/codes/linux/kernel/softirq.c:472)
do_softirq() (/data00/codes/linux/kernel/softirq.c:459)

```

调试kernel收发包流程

用 `curl http://bilibili.com` 测试，在 `__netif_receive_skb_core` 下断点

bt如下

```

Breakpoint 2, __netif_receive_skb_core
(pskb=pskb@entry=0xfffffc900000003da0,
  pfmemalloc=pfmemalloc@entry=false,
  ppt_prev=ppt_prev@entry=0xfffffc900000003da8)
  at net/core/dev.c:5240
5240    static int __netif_receive_skb_core(struct sk_buff
**pskb, bool pfmemalloc,
(gdb) frame
#0  __netif_receive_skb_core
(pskb=pskb@entry=0xfffffc900000003da0,
  pfmemalloc=pfmemalloc@entry=false,
  ppt_prev=ppt_prev@entry=0xfffffc900000003da8)
  at net/core/dev.c:5240
5240    static int __netif_receive_skb_core(struct sk_buff
**pskb, bool pfmemalloc,
(gdb) bt
#0  __netif_receive_skb_core
(pskb=pskb@entry=0xfffffc900000003da0,
  pfmemalloc=pfmemalloc@entry=false,
  ppt_prev=ppt_prev@entry=0xfffffc900000003da8)
  at net/core/dev.c:5240
#1  0xffffffff81bac595 in __netif_receive_skb_list_core (
  head=head@entry=0xfffff888003ddccb8,
  pfmemalloc=pfmemalloc@entry=false)
  at net/core/dev.c:5528
#2  0xffffffff81bacd3b in __netif_receive_skb_list
(head=0xfffff888003ddccb8)
  at net/core/dev.c:5595
#3  netif_receive_skb_list_internal
(head=head@entry=0xfffff888003ddccb8)
  at net/core/dev.c:5686
#4  0xffffffff81bad4be in gro_normal_list
(napi=0xfffff888003ddcbb0)
  at ./include/net/gro.h:439
#5  gro_normal_list (napi=0xfffff888003ddcbb0) at
./include/net/gro.h:435
#6  napi_complete_done (n=n@entry=0xfffff888003ddcbb0,
work_done=<optimized out>)
  at net/core/dev.c:6026

```

```
#7  0xffffffff819b8863 in e1000_clean (napi=0xffff888003ddcbb0,
    budget=64)
    at drivers/net/ethernet/intel/e1000/e1000_main.c:3811
#8  0xffffffff81bad627 in __napi_poll
    (n=n@entry=0xffff888003ddcbb0,
    repoll=repoll@entry=0xffffc90000003f47) at
    net/core/dev.c:6460
#9  0xffffffff81badcb2 in napi_poll (repoll=0xffffc90000003f58,
    n=0xffff888003ddcbb0)
    at net/core/dev.c:6527
#10 net_rx_action (h=<optimized out>) at net/core/dev.c:6660
#11 0xffffffff81f0a5a0 in __do_softirq () at
    kernel/softirq.c:553
#12 0xffffffff81091a8e in do_softirq () at kernel/softirq.c:454
#13 do_softirq () at kernel/softirq.c:441
```

[inet_init](#) 添加 ipv4 的 packet handler

receive_skb* 有list和非list版本，list可以做GRO合并，根据条件将多个skb合并，减少重复处理次数

调用顺序

```
netif\_receive\_skb\_list
__netif\_receive\_skb\_list\_core
__netif\_receive\_skb\_core
deliver\_skb
```

回调packet_type的handler
5.

路由选择

[__fib_lookup](#)

[fib_lookup](#)

struct rtable

代表一条路由规则，路由结束绑定到[skb#_skb_refdst](#)

- struct dst_entry dst;
 - 下面详细介绍。
- _u16 rt_type
 - 本路由类型

字段	值	ip route	含义
RTN_UNICAST	1	unicast	默认值，如果skb#_skb_refdst是此值，数据需要被转发
RTN_LOCAL	2	local	本机路由
RNT_BROADCAST	3	广播	需要转发
RTN_MULTICAST	5	多播	多播
RNT_UNREACHABLE	7	unreachable	丢弃，返回ICMP network unreachable

- __u8 rt_uses_gateway
 - bool，路由的下一跳是网关(ip route 返回类似 **via 10.0.0.1** 的格式)，那么 **rt_gw4** 包含网关IP地址。
- u8 rt_gw_family
 - 如果rt_uses_gateway是0，那rt_gw_family是0。如果网关地址是IPV4，=AF_INET。IPV6，=AF_INET6
- union {__be32 rt_gw4; struct in6_addr rt_gw6;}
 - 如果是网关，根据IP类型使用 **rt_gw4** 或 **rt_gw6** 字段

struct dst_entry

- struct net_device *dev
 - 发送数据的网络设备，数据包最终从此设备发送
- struct xfrm_state *xfrm
 - 和IPsec 相关，一般是NULL
- int (_input)(struct sk_buff *skb)
 - 根据路由结果，input是不同值。决定数据包接下来如果处理。dest ip本机则继续向上传递。不是本机的ip则转发

可选值	备注
dst_discard	默认值
ip_local_deliver	目的地址是本机，或者广播数据包
ip_forward	单播，但不是本机，需要转发
ip_error	没有找到匹配路由， unreachable ，数据包丢弃，返回ICMP host unreachable

可选值	备注
lwtunnel_input	Kconfig - net/Kconfig - Linux source code (v5.14.7) - Bootlin

- int (_output)(struct net_net, struct sock _sk, struct sk_buff_skb);
 - 根据路由结果output指向不同函数。决定数据包如何发送

可选值	备注
dst_discard_out()	默认值
ip_output	本机生成，需要发送的单播包
ip_rt_bug()	bug?
xfrm4_output	要转发的数据包
ip_mc_output	本机生成，需发送的多播（multicast）包

skb 到 sock

UDP sock send to skb

✓ 3: tid=3

PAUSED ON BREAKPOINT

__skb_insert	skbuff.h	2200:2
__skb_queue_before	skbuff.h	2309:2
__skb_queue_tail	skbuff.h	2342:2
__udp_enqueue_schedule_skb	udp.c	1537:2
__udp_queue_rcv_skb	udp.c	2037:7
udp_queue_rcv_one_skb	udp.c	2166:9
udp_queue_rcv_skb	udp.c	2184:10
udp_unicast_rcv_skb	udp.c	2344:8
__udp4_lib_rcv	udp.c	2420:10
udp_rcv	udp.c	2602:9
ip_protocol_deliver_rcu	ip_input.c	205:9
ip_local_deliver_finish	ip_input.c	233:2
NF_HOOK	netfilter.h	304:9
ip_local_deliver	ip_input.c	254:9
dst_input	dst.h	468:9
ip_sublist_rcv_finish	ip_input.c	580:3
ip_list_rcv_finish	ip_input.c	631:2
ip_sublist_rcv	ip_input.c	639:2
ip_list_rcv	ip_input.c	674:3
__netif_receive_skb_list_ptype	dev.c	5570:3

[Load More Stack Frames](#)

> 4: tid=4

PAUSED

tcp sock send

__skb_insert	skbuff.h	2200:2
__skb_queue_before	skbuff.h	2309:2
__skb_queue_tail	skbuff.h	2342:2
tcp_add_write_queue_tail	tcp.h	1942:2
tcp_skb_entail	tcp.c	670:2
tcp_sendmsg_locked	tcp.c	1157:4
tcp_sendmsg	tcp.c	1336:8
inet_sendmsg	af_inet.c	840:9
sock_sendmsg_nosec	socket.c	730:12
__sock_sendmsg	socket.c	745:16
__sys_sendto	socket.c	2194:8
__do_sys_sendto	socket.c	2206:9
__se_sys_sendto	socket.c	2202:1
__x64_sys_sendto	socket.c	2202:1
do_syscall_x64	common.c	50:14
do_syscall_64	common.c	80:7
entry_SYSCALL_64	entry_64.S	120
4B	@4b..8b	3
> 4: tid=4		
<div>PAUSED</div> <div>     </div>		

tcp sock receive from skb

✓ 3: tid=3

PAUSED ON BREAKPOINT

__skb_insert	skbuff.h	2200:2
__skb_queue_before	skbuff.h	2309:2
__skb_queue_tail	skbuff.h	2342:2
tcp_queue_rcv	tcp_input.c	4973:3
tcp_data_queue	tcp_input.c	5090:11
tcp_rcv_established	tcp_input.c	6053:2
tcp_v4_do_rcv	tcp_ipv4.c	1728:3
tcp_v4_rcv	tcp_ipv4.c	2150:9
ip_protocol_deliver_rcu	ip_input.c	205:9
ip_local_deliver_finish	ip_input.c	233:2
NF_HOOK	netfilter.h	304:9
ip_local_deliver	ip_input.c	254:9
dst_input	dst.h	468:9
ip_sublist_rcv_finish	ip_input.c	580:3
ip_list_rcv_finish	ip_input.c	631:2
ip_sublist_rcv	ip_input.c	639:2
ip_list_rcv	ip_input.c	674:3
__netif_receive_skb_list_ptype	dev.c	5570:3
__netif_receive_skb_list_core	dev.c	5618:2
__netif_receive_skb_list	dev.c	5670:3

[Load More Stack Frames](#)

参考

[Routing Decisions in the Linux Kernel - Part 1: Lookup and packet flow](#)
[\[Thermalcircle.de\]](#)
[networking:kernel_flow \[Wiki\]](#)